

De Novo Assembly of High-throughput Short Read Sequences

Chuming Chen

Center for Bioinformatics and Computational Biology (CBCB)
University of Delaware

NECC Third Skate Genome Annotation Workshop
May 23, 2011

Outline

- Genome assembly primer
- High-throughput short read sequencing (NGS) assembly pipeline
- Case study

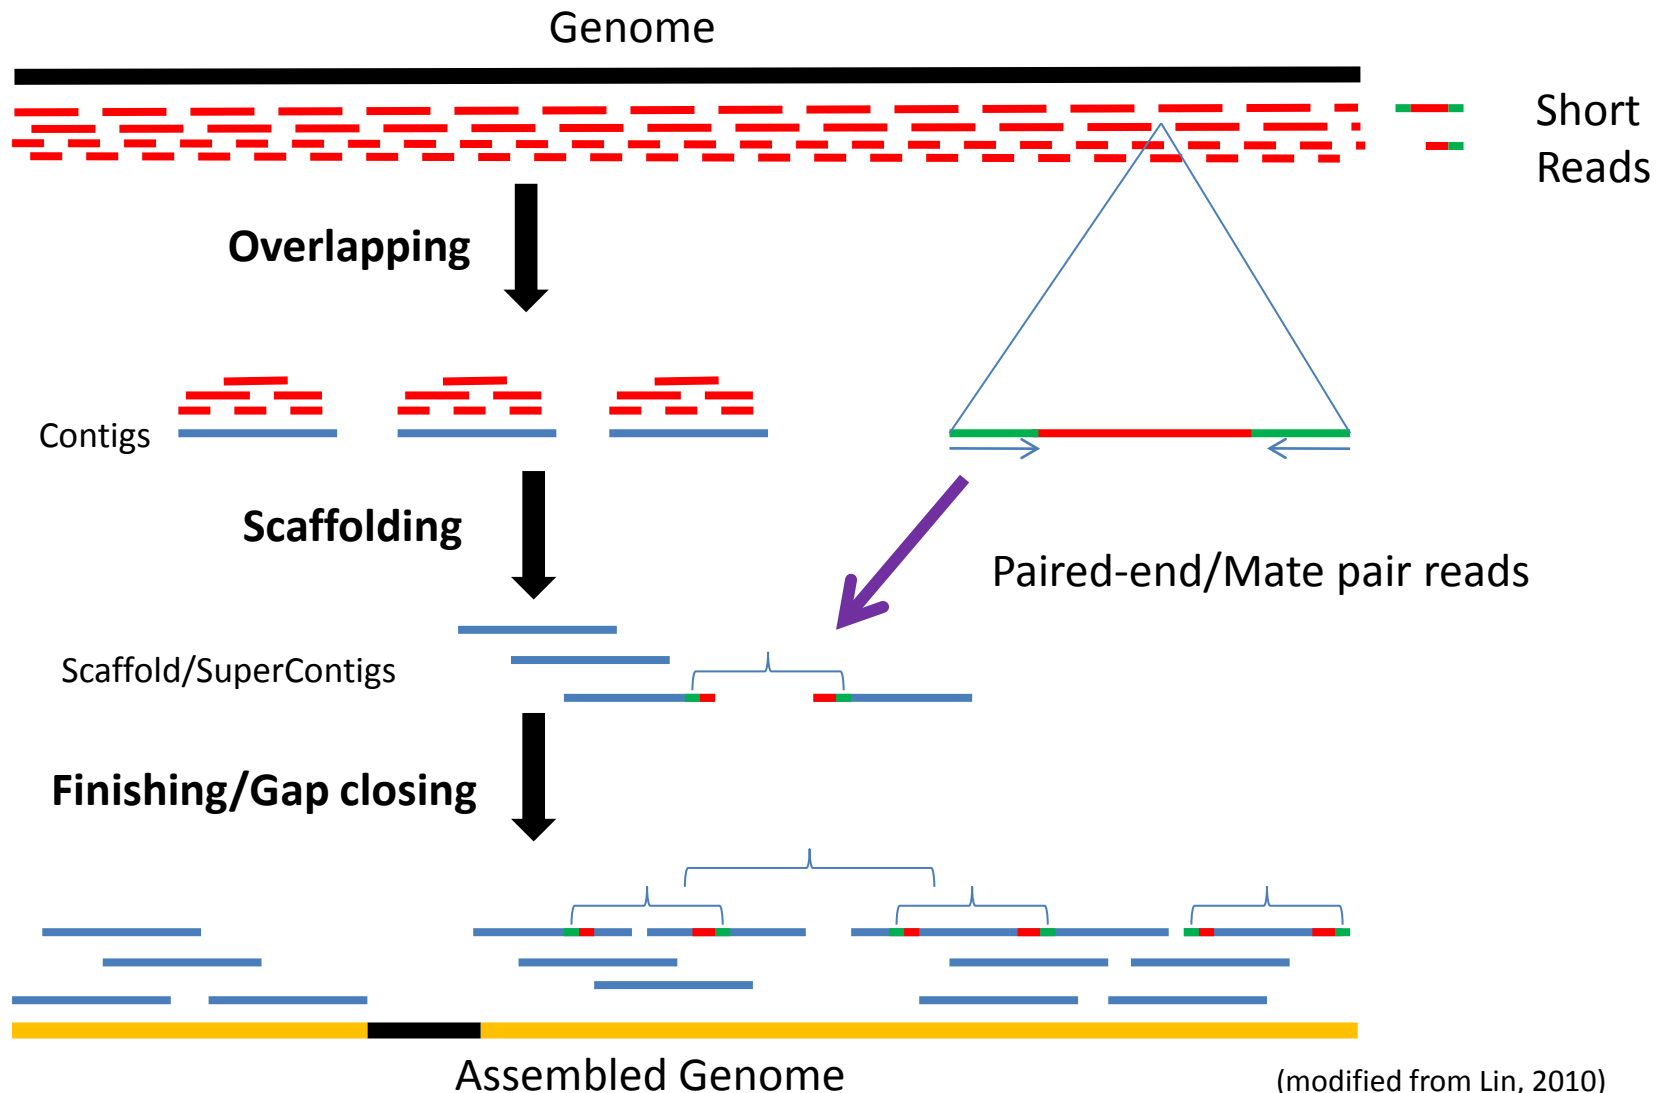
Genome sequencing and assembly

- Genome
 - Long stretches of contiguous DNA sequences (base pairs)
 - Different genome sizes (i. e. virus: 3.5k, Human: 3.3 billion)
- Genome sequencers (NHGRI, Feb. 4, 2011)
 - Sanger-based sequencing (500-600 bases)
 - 454 sequencing (300-400 bases)
 - Illumina and SOLiD sequencing (50-100 bases)
- Sequencing and assembly
 - A genome must be fragmented, sequenced piece by piece and then re-assembled to obtain the full contiguous sequence


Assembly approaches

- Hierarchical assembly
 - Mapping the genome to a set of large insert clones
 - Reduce the assembly of the sequencing reads from the entire genome to a single clone, typically 40 - 200 Kb
 - The genome sequence is then assembled by aligning sequences of adjacent clones
- Whole genome shot-gun assembly
 - The entire genome is fragmented
 - The shotgun process takes reads from random positions along the chromosomes that make up one genome
 - The assembler then reconstructs the reads up to the chromosome length
 - Assembly is possible because the target is over-sampled by the shotgun reads, such that reads overlap


Whole genome shot-gun sequencing and assembly







Paired-end vs. mate-pair reads



(Fass, 2010)

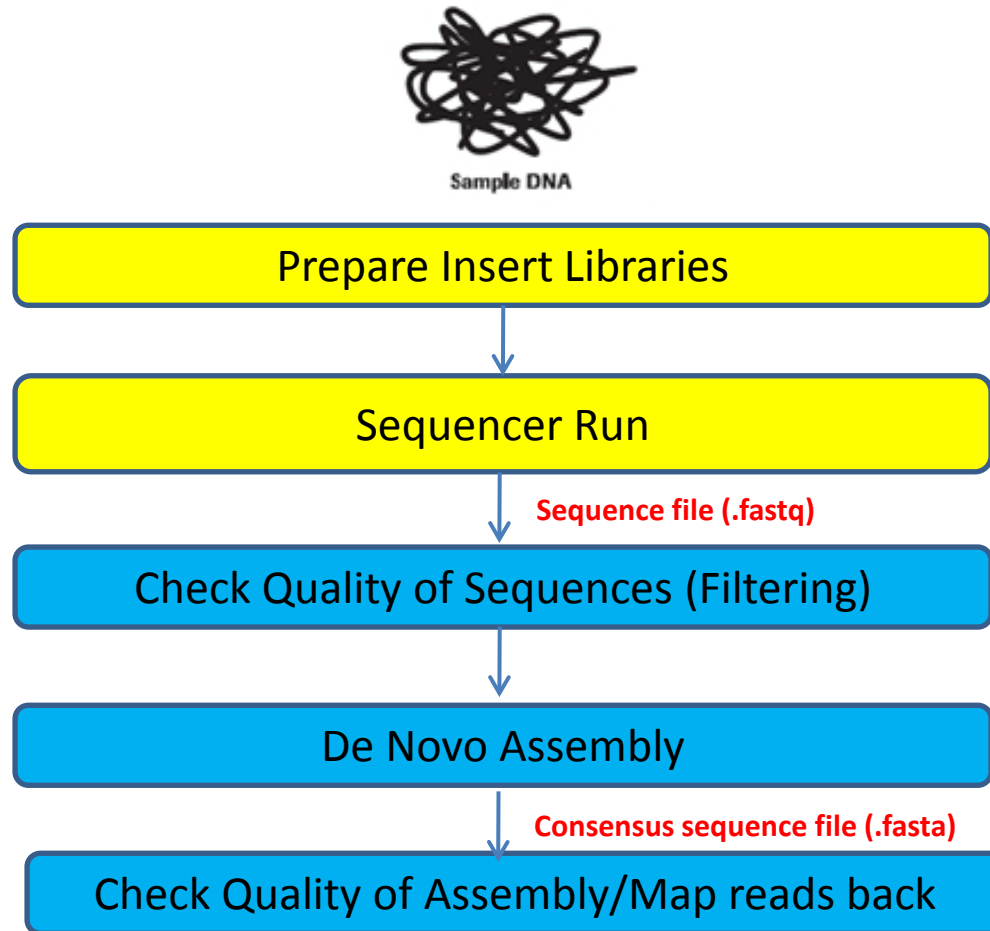


(Fass, 2010)

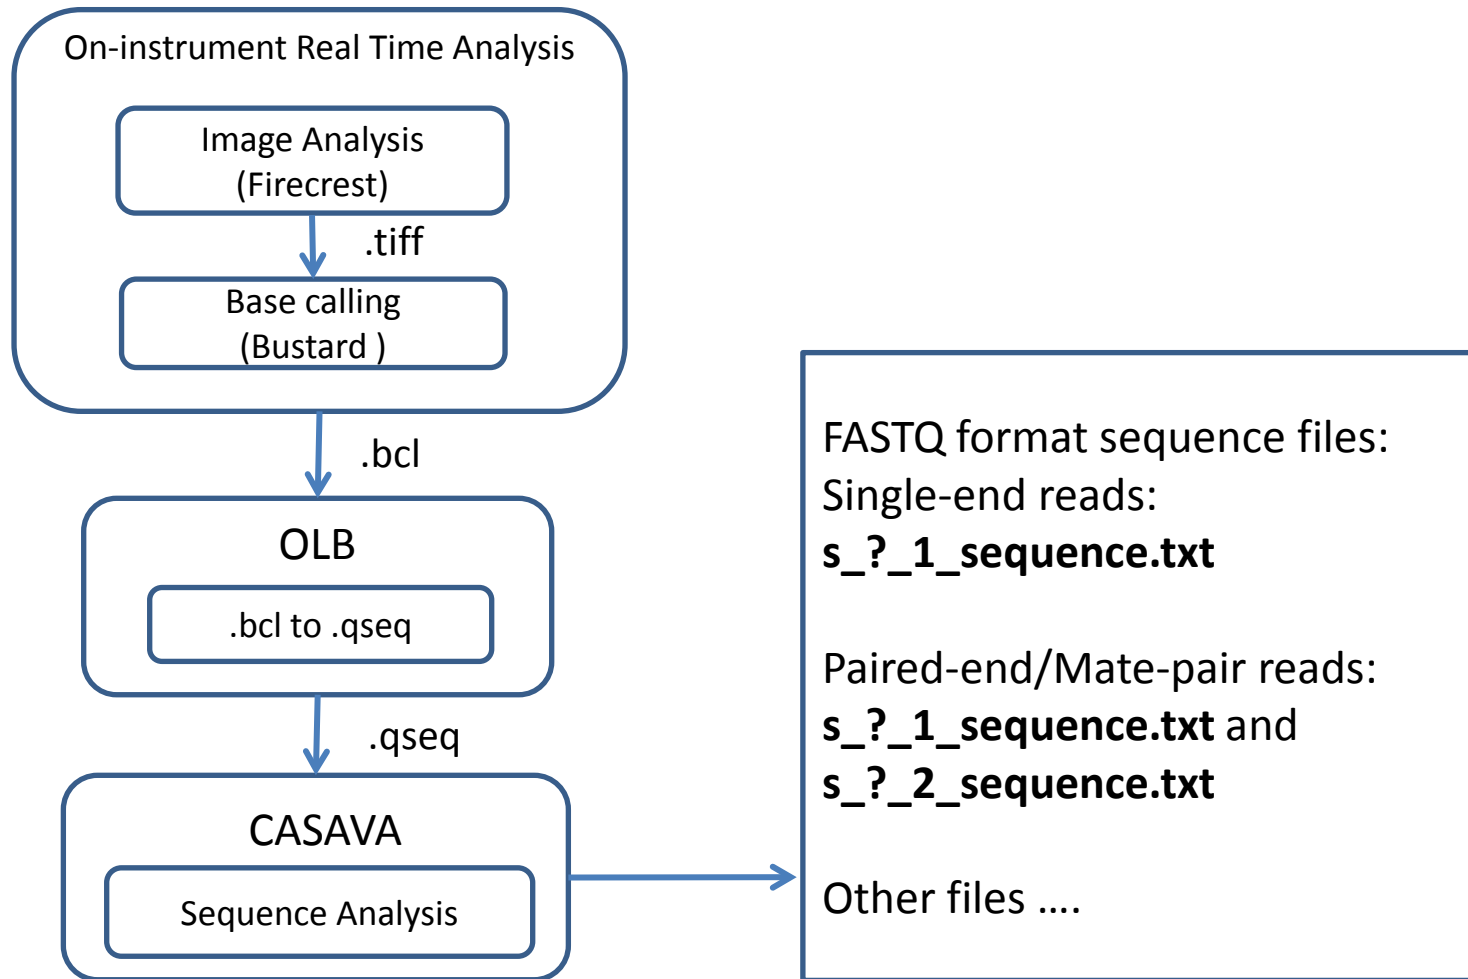


(Phillippy, 2008)

Assembly pipeline



HiSeq2000 sequence data processing pipeline



FASTQ format sequence files

s_?_1_sequence.txt

The diagram illustrates the FASTQ format structure with four labels and arrows pointing to the corresponding parts of the first sequence record in s_?_1_sequence.txt:

- @Sequence ID**: Points to the first line of the record: `@HWI-ST741_0085:1:1101:1444:1939#0/1`
- Sequence (A, C, G, T, N)**: Points to the second line of the record: `ATAGTTACAATCGATCCATTGTCAGAGTACAGATACATGATACGGGAATACCATTTCATGGGAAAGCCAGCAAAGTTCGATCAAGGATAGACCGAGG`
- +Sequence ID**: Points to the third line of the record: `+HWI-ST741_0085:1:1101:1444:1939#0/1`
- ASCII encoded quality score**: Points to the fourth line of the record: `fffffdfffffgggfafffcdfcffffbddddeaegfgfgafaffW^a]a`b^d[dbdbbbcccccb[`dW\a`a`^caa[ffcffafaccfcd_fcfc_cb`

...
@HWI-ST741_0085:1:1101:1417:1939#0/1
GCTCAGTACTGGGACCCCAGCTATTTACAAAATATATTAATGATCTGGATGAGGGATTTGAAGGCAATATCTCCAAGTTTGC GGATGACATTAAGCTGGGG
+HWI-ST741_0085:1:1101:1417:1939#0/1
hhhhghhhhhhhhhhhhhhhhhhhhhhehhhhhhhhghhhhhhhhhhhhhghhhghhghhhhhahfghhhfhhhhhggafffhhhhdghg[hghhfhhhhageh
...

s_?_2_sequence.txt

...
@HWI-ST741_0085:1:1101:1444:1939#0/2
CCCAGCTTATCCTTGCAACTCTTCTTAAATAGAGGCACAACATTAATCACCTCCCTTCTATGGACTCCTTTTATACCTCACCTTCCTCGGCAGGGCCAG
+HWI-ST741_0085:1:1101:1444:1939#0/2
edeaadffffcaffcdaeaeffdfdecfefaceccfdffdfddfffffdfffdcaffffcfffagggggggg`ggggggfggggggggd]eedaa_`BB
@HWI-ST741_0085:1:1101:1417:1939#0/2
ACCATTGCTTTCTTTACTGCCTGCTGCACCTGCATGCCTACCTTCAATGACTGGTGTACCATGACACCCAGGTCTCGCTGCATCTCCCCCTTTCCCAATC
+HWI-ST741_0085:1:1101:1417:1939#0/2
hhhhghhhhhhhhhhhhhhhhhhhhhghhhhdhhghhhhhhhhhhhchhhhhhhhhhhfghhhhhfhhdhfhhddhahhhfefhhfdfffedhhhehhhhhehhcf f f f f d
...

A closer look at sequence file

1	2	3	4	5	6	7
@HWI-ST741_0085:1:1101:1444:1939#0/1						
ATAGTTACAATCGATCCATTTGCAGAGTACAGATACATGATACGGAAT						
+HWI-ST741_0085:1:1101:1444:1939#0/1						
ffffdffdfffffgggfafffcdfcffbfdddeaegfgfgafaffw^a]						

1	2	3	4	5	6	7
@HWI-ST741_0085:1:1101:1444:1939#0/2						
CCCAGCTTATCCTTGCAACTCTTCTTAAATAGAGGCACAACATTAATCA						
+HWI-ST741_0085:1:1101:1444:1939#0/2						
Edeaadffffcaffcdaeaefdfdecfefaceccfdffdfddfffffd						

1. the unique instrument name
2. flowcell lane
3. tile number within the flowcell lane
4. 'x'-coordinate of the cluster within the tile
5. 'y'-coordinate of the cluster within the tile
6. index number for a multiplexed sample (0 for no multiplexing)
7. the member of a pair, /1 or /2 (*paired-end or mate-pair reads only*)

Quality scores

- Phred quality scores Q are defined as a property which is logarithmically related to the base-calling error probabilities P
- A Phred score of a base is: $Q_{\text{phred}} = -10 \log_{10} P$, where P is the estimated probability of a base being wrong

Phred Quality Score	Probability of incorrect base call	Base call accuracy
10	1 in 10	90 %
20	1 in 100	99 %
30	1 in 1000	99.9 %
40	1 in 10000	99.99 %
50	1 in 100000	99.999 %

- Phred scoring scheme, encoded as an ASCII character by adding 64 to the Phred value


```
@HWI-ST741_0085:1:1101:1444:1939#0/1
ATAGTTACAATCGATCCATTTCAGAGTACAGATACATGATACGGGAAT
+HWI-ST741_0085:1:1101:1444:1939#0/1
ffffdfffffgggfaaffcfdcfbfdddeaegfgfgafaffw^a] HiSeq score
                                     ↓
39 33 38 33 38 38 23 30 33 29 Phred score
```

(Wikipedia.org)

Sequence reads quality assessment

- **FastQC** (Baraham Bioinformatics, UK)
 - Basic Statistics
 - Per Base Sequence Quality
 - Per Sequence Quality Scores
 - Per Base Sequence Content
 - Per Base GC Content
 - Per Sequence GC Content
 - Per Base N Content
 - Sequence Length Distribution
 - Duplicate Sequences
 - Overrepresented Sequences
 - Overrepresented Kmers

Per base sequence quality



Trim Sequences

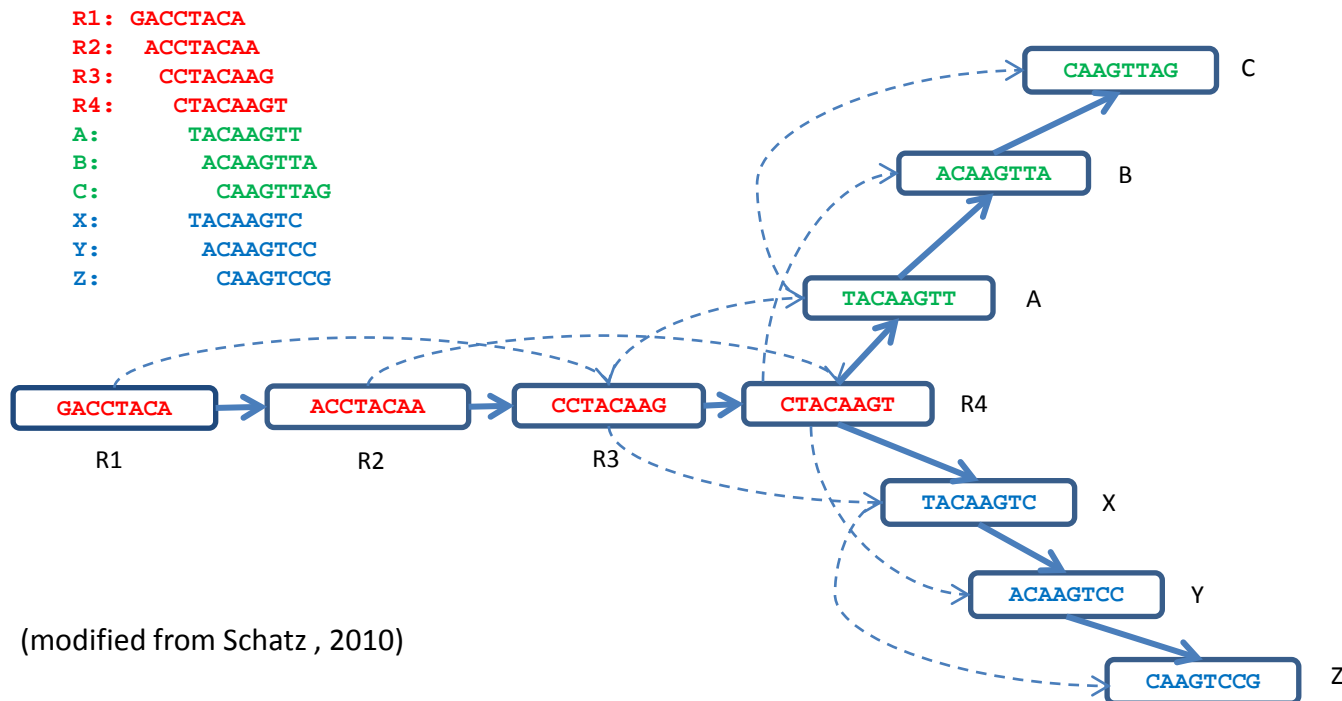
- Quality trimming
 - Based on quality scores
- Ambiguity trimming
 - Remove stretches of Ns
- Adapter sequence trimming
 - Remove sequence adapters
- Base trim
 - Remove a specified number of bases at either 3' or 5' end of the reads
- Length trimming
 - Remove reads shorter or longer than a specified threshold

De Novo Assembly algorithms

- Overlap/Layout/Consensus Graph
- de Bruijn Graph


Overlap/Layout/Consensus graphs

- A node corresponds to a read, an edge denotes an overlap between two reads.
- The overlap graph is used to compute a layout of reads and consensus sequence of contigs by pair-wise sequence alignment.
- Good for sequences with limited number of reads but significant overlap. Computational intensive for short reads (short and high error rate).
- Example assemblers: Celera Assembler, Arachne, CAP and PCAP



de Bruijn graphs

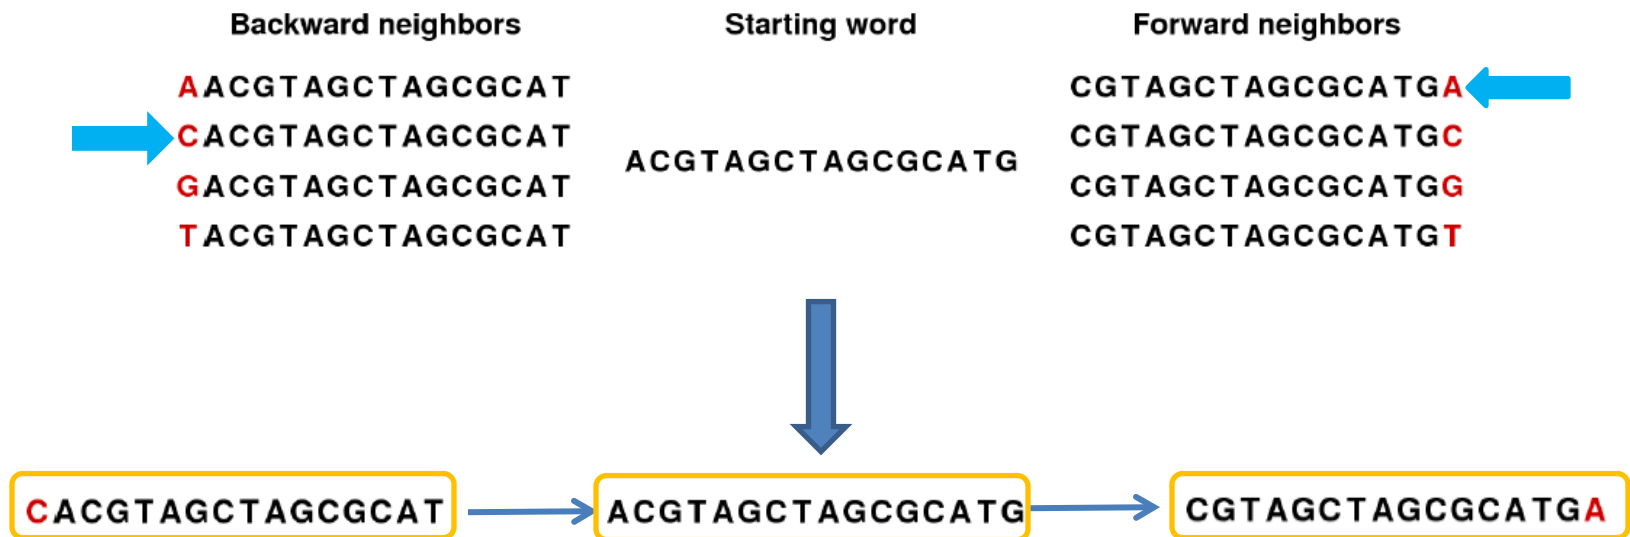
- No need for all against all overlap discovery.
- Break reads into smaller sequences of DNA (K-mers, K denotes the length in bases of these sequences).
- Captures overlaps of length K-1 between these K-mers.
- More sensitive to repeats and sequencing errors.
- By construction, the graph contains a path corresponding to the original sequence.
- Example assemblers: Euler, Velvet, ABySS, AllPaths, SOAPdenovo, CLC Bio



(modified from Schatz , 2010)

CLC Bio De Novo assembly

- Make a table of the words (K-mers) seen in the reads.
- Build de Bruijn graph from the word table.
- Use the reads to resolve the repeats.
- Use the information from paired reads to resolve larger repeats.
- Output resulting contigs based on the paths.



(modified from CLC Bio, 2011)

Word (K-mer) size

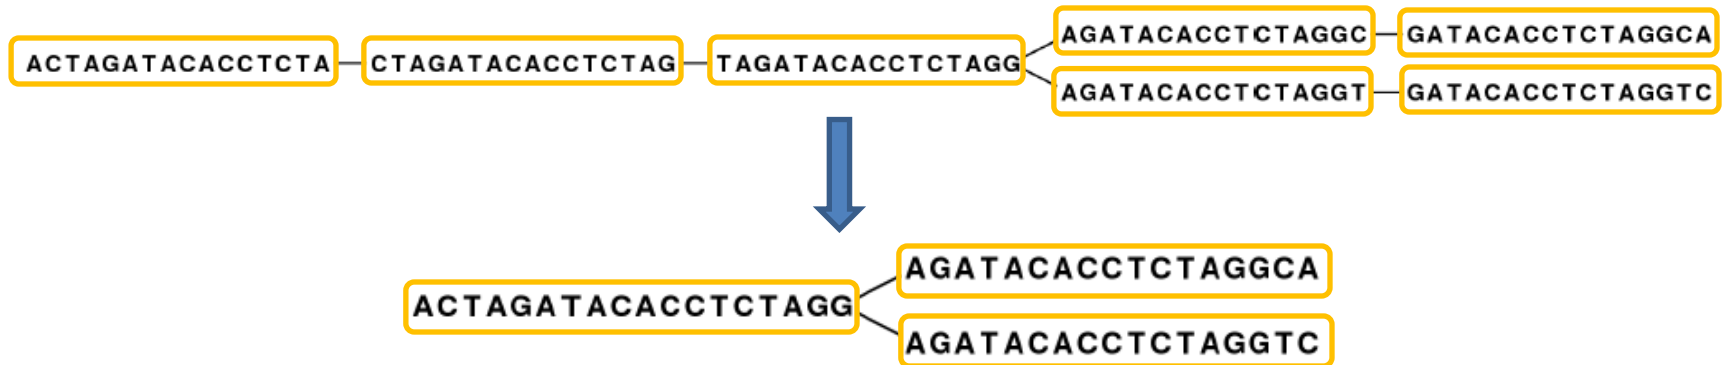
To strike a balance, CLC bio's de novo assembler chooses a word length based on the amount of input data: the more data, the longer the word length. It is based on the following:

word size 12: 0 bp - 30000 bp
word size 13: 30001 bp - 90002 bp
word size 14: 90003 bp - 270008 bp
word size 15: 270009 bp - 810026 bp
word size 16: 810027 bp - 2430080 bp
word size 17: 2430081 bp - 7290242 bp
word size 18: 7290243 bp - 21870728 bp
word size 19: 21870729 bp - 65612186 bp
word size 20: 65612187 bp - 196836560 bp
word size 21: 196836561 bp - 590509682 bp
word size 22: 590509683 bp - 1771529048 bp
word size 23: 1771529049 bp - 5314587146 bp
word size 24: 5314587147 bp - 15943761440 bp
word size 25: 15943761441 bp - 47831284322 bp
word size 26: 47831284323 bp - 143493852968 bp
word size 27: 143493852969 bp - 430481558906 bp
word size 28: 430481558907 bp - 1291444676720 bp
word size 29: 1291444676721 bp - 3874334030162 bp
word size 30: 3874334030163 bp - 11623002090488 bp
word size 31: 11623002090489 bp and up

(CLC Bio, 2011)

Repeats or sequencing errors

Graph Reduction



SNP or Sequencing Error



Repeat Sequence



(modified from CLC Bio, 2011)

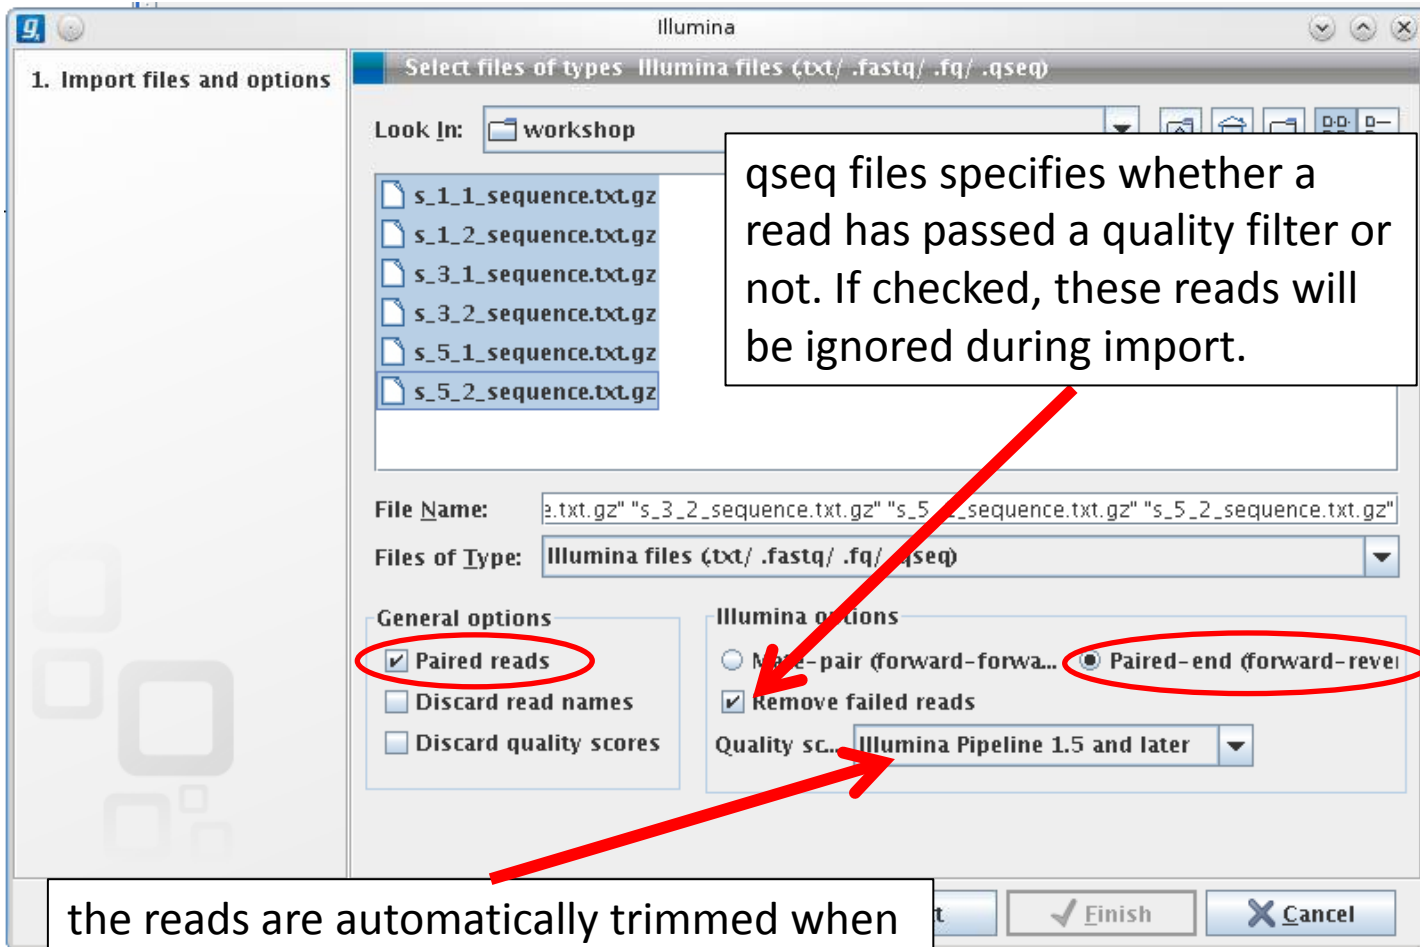
Assembly quality assessment

- Continuity
 - Lengths distribution of contigs/scaffolds.
 - Average length, minimum and maximum lengths, combined total lengths.
 - **N50** captures how much of the assembly is covered by relatively large contigs.
 - The N50 is the length of the smallest contig in the set that contains the fewest (largest) contigs whose combined length represents at least 50% of the assembly.
 - Compute N50
 - first ordering all contigs (or scaffolds) by length,
 - then summing up their lengths (starting with the longest) until the sum exceeds 50% of the total length of all contigs.
 - the corresponding contig length is N50.
- Accuracy or “Correctness”
 - **Base accuracy** – the frequency of calling the correct nucleotide at a given position in the assembly.
 - **Mis-assembly rate** – the frequency of rearrangements, significant insertions, deletions and inversions.

Case study

- Show the basic steps involved in De Novo assembly of high-throughput short read sequences
- Data
 - 3 lanes of Illumina HiSeq 2000 short read sequences for the little skate (a couple weeks ago)
- Assembler
 - CLC Bio Genomics Workbench 4.6

Import sequence data



qseq files specifies whether a read has passed a quality filter or not. If checked, these reads will be ignored during import.

the reads are automatically trimmed when a quality score B is encountered in the input file

Quality trimming

1. Convert quality score (Q) to error probability: $P_{\text{error}} = 10^{Q/-10}$, low values are high quality bases.
2. For every base a new value is calculated: Limit – P_{error} , negative for low quality bases.
3. For every base, the running sum of this value is calculated. If the sum drops below zero, it is set to zero.
4. The region retained is between the first positive value of the running sum and the highest value of the running sum. Everything before and after this region will be trimmed off.

1. Select sequencing data

2. Quality trimming

Set parameters

Quality trimming

☒ Trim using quality scores

Limit: 0.05

☒ Trim ambiguous nucleotides

Maximum number of ambiguities: 2

? [Help Icon] Previous Next Finish Cancel

If this maximum is set to e.g. 2, the algorithm finds the maximum length region containing 2 or fewer ambiguities and then trims away the ends not included in this region.

Adapter trimming

1. Remove adapter
2. Discard when not found
3. Discard when found

The screenshot displays the 'Trim Sequences' application window. On the left, a sidebar lists three steps: '1. Select sequencing data', '2. Quality trimming', and '3. Adapter trimming'. The main window is titled 'Trim Sequences' and has a 'Set parameters' tab. Under 'Adapter trimming', the 'Search on both strands' checkbox is checked and circled in red. Below this is a table listing various adapters. The 'Illumina HiSeq 2000' adapter is selected, and its 'Action' is 'Remove adapter'. A red arrow points from the list to the 'Edit Alignment scores' dialog box. This dialog box shows 'Costs' with 'Mismatch cost' set to 2 and 'Gap cost' set to 3. Under 'Match thresholds', 'Allow internal matches' and 'Allow end matches' are both checked, with 'Minimum score' set to 10 and 'Minimum score at end' set to 4. At the bottom of the main window, a 'Preview' section shows 'Number of reads' as 1,000 and 'Number of nucleotides' as 92. A table below this shows 'Illumina HiSeq 2000 ...' with 19 matches found and 0 reads discarded. Navigation buttons like '?', 'Previous', and 'Next' are at the bottom.

Adapter trimming parameters:

Name	Sequence	Strand	Alignment score	Action
5' adapter small RNA	ATCGTAGGCACCT...	Minus	3, 5, 15, 2	Remove adapter
454 Sequence Prim...	GCCTCCCTCGCGC...	Plus	3, 2, 15, 2	Remove adapter
454 Sequence Prim...	GCCTTGCCAGCCC...	Minus	3, 2, 15, 2	Remove adapter
454 miRNA forwar...	GCCTCCCTCGCGC...	Plus	3, 2, 15, 2	Discard when not f...
454 miRNA revers...	GCCTTGCCAGCCC...	Minus	3, 2, 15, 2	Discard when not f...
<input checked="" type="checkbox"/> Illumina HiSeq 200...	ACACTCTTTCCTA...	Plus	2, 3, ... 4	Remove adapter

Preview:

Number of reads: 1,000 Number of nucleotides: 92

Name	Matches found	Reads discarded
Illumina HiSeq 2000 ...	19	0

Edit Alignment scores dialog:

Costs

Mismatch cost: 2
Gap cost: 3

Match thresholds

☒ Allow internal matches
Minimum score: 10
☒ Allow end matches
Minimum score at end: 4

Sequence filtering

The screenshot shows a software window titled "Trim Sequences". On the left is a sidebar with a list of steps: "1. Select sequencing data", "2. Quality trimming", "3. Adapter trimming", and "4. Sequence filtering". The main area is titled "Set parameters" and contains two sections. The "Trim bases" section has two checkboxes: "Remove 5' terminal nucleotides" and "Remove 3' terminal nucleotides", each followed by a numeric input field set to "1". The "Filter on length" section has two checkboxes: "Discard reads below leng..." and "Discard reads above length", followed by numeric input fields set to "15" and "1,000" respectively. At the bottom are buttons for "?", a circular arrow, "Previous", "Next", "Finish", and "Cancel".

Trim Sequences

1. Select sequencing data
2. Quality trimming
3. Adapter trimming
4. Sequence filtering

Set parameters

Trim bases

☐ Remove 5' terminal nucleotides

☐ Remove 3' terminal nucleotides

Filter on length

☐ Discard reads below leng...

☐ Discard reads above length

? ↻ ⬅ Previous ➡ Next ✓ Finish ✕ Cancel

Trimming result handling

The screenshot shows a software window titled "Trim Sequences". On the left is a sidebar with a list of five steps: "1. Select sequencing data", "2. Quality trimming", "3. Adapter trimming", "4. Sequence filtering", and "5. Result handling". The "5. Result handling" step is highlighted. The main area of the window is titled "Result handling" and contains three sections: "Output options" with two checked checkboxes, "Create list of discarded sequences" and "Create report"; "Result handling" with two radio buttons, "Open" and "Save", where "Save" is selected; and "Log handling" with one checked checkbox, "Make log". At the bottom of the window are four buttons: a help button with a question mark, a back button with a left arrow and the text "Previous", a forward button with a right arrow and the text "Next", an "Finish" button with a checkmark, and a "Cancel" button with an X.

Trim Sequences

1. Select sequencing data
2. Quality trimming
3. Adapter trimming
4. Sequence filtering
5. Result handling

Result handling

Output options

- ☒ Create list of discarded sequences
- ☒ Create report

Result handling

- ☐ Open
- ☒ Save

Log handling

- ☒ Make log

? ↶ Previous ↷ Next Finish Cancel

Trimming report

Summary

Name	Number of reads	Avg. length	Number of reads after trim	Percentage retained	Avg. length after trim
S_1_1_sequence (paired)	190,790,624	93.4	190,784,343	~100%	93.1
S_1_1_sequence	1,655,955	63.9	1,654,977	99.94%	63.0
S_3_1_sequence (paired)	209,140,424	92.9	209,131,515	~100%	92.6
S_3_1_sequence	2,016,294	62.4	2,015,199	99.95%	61.8
S_5_1_sequence (paired)	223,212,034	92.4	223,201,524	~100%	92.1
S_5_1_sequence	2,373,091	62.2	2,371,693	99.98%	61.5

Trim settings

- Removal of low quality sequence. (limit = 0.05).
- Removal of ambiguous nucleotides: maximal 2 nucleotides allowed.
- Removal of adapter sequences, using the following adapters :
Illumina HiSeq 2000 PE Adapter (ACACTCTTTCCCTACACGACGTCTTCCGATCT), strand = Plus, action = Remove adapter, score = [2, 3, 10, 4]

Detailed results

Trim	Input reads	No trim	Trimmed	Nothing left or Discarded
Trim on quality	629,187,422	618,707,269	10,474,178	5,975
Ambiguity trim	629,181,447	628,910,632	256,018	14,797
Adapter trimming	629,166,650	617,589,894	11,569,357	7,399

Assembly parameters

De Novo Assembly

1. Select sequencing reads
2. Set parameters

Set parameters

Selected reads

	Input	Length	Type	Settings
	s_1_1_sequence (paired) ...	Long	Paired	Max distance=500
	s_1_1_sequence (paired) ...	Long	Single	Default
	s_1_1_sequence trimmed	Long	Single	Default
	s_3_1_sequence (paired) ...	Long	Paired	Max distance=500
	s_3_1_sequence (paired) ...	Long	Single	Default

Long reads mapping parameters

Mismatch cost

Insertion cost

Deletion cost

Length fraction

Similarity

☐ Global alignment

☐ Color space alignment

Colorspace error cost

Paired parameters

Minimum distance

Maximum distance

☐ Guidance only

Minimum length fraction of reads must match the reference

Minimum fraction of similarity between read and the reference

If checked, end gaps will be treated as mismatches

? ↶ ↷ ↵ ↶ ↷

Previous Next Finish Cancel

General assembly options

De Novo Assembly

1. Select sequencing reads
2. Set parameters
3. Set general assembly options

Set parameters

Options

☒ Add conflict annotations

Conflict resolution

☐ Unknown nucleotide (N)
☐ Ambiguity nucleotides (R, Y, etc.)
☒ Vote (A, C, G, T)

Non-specific matches

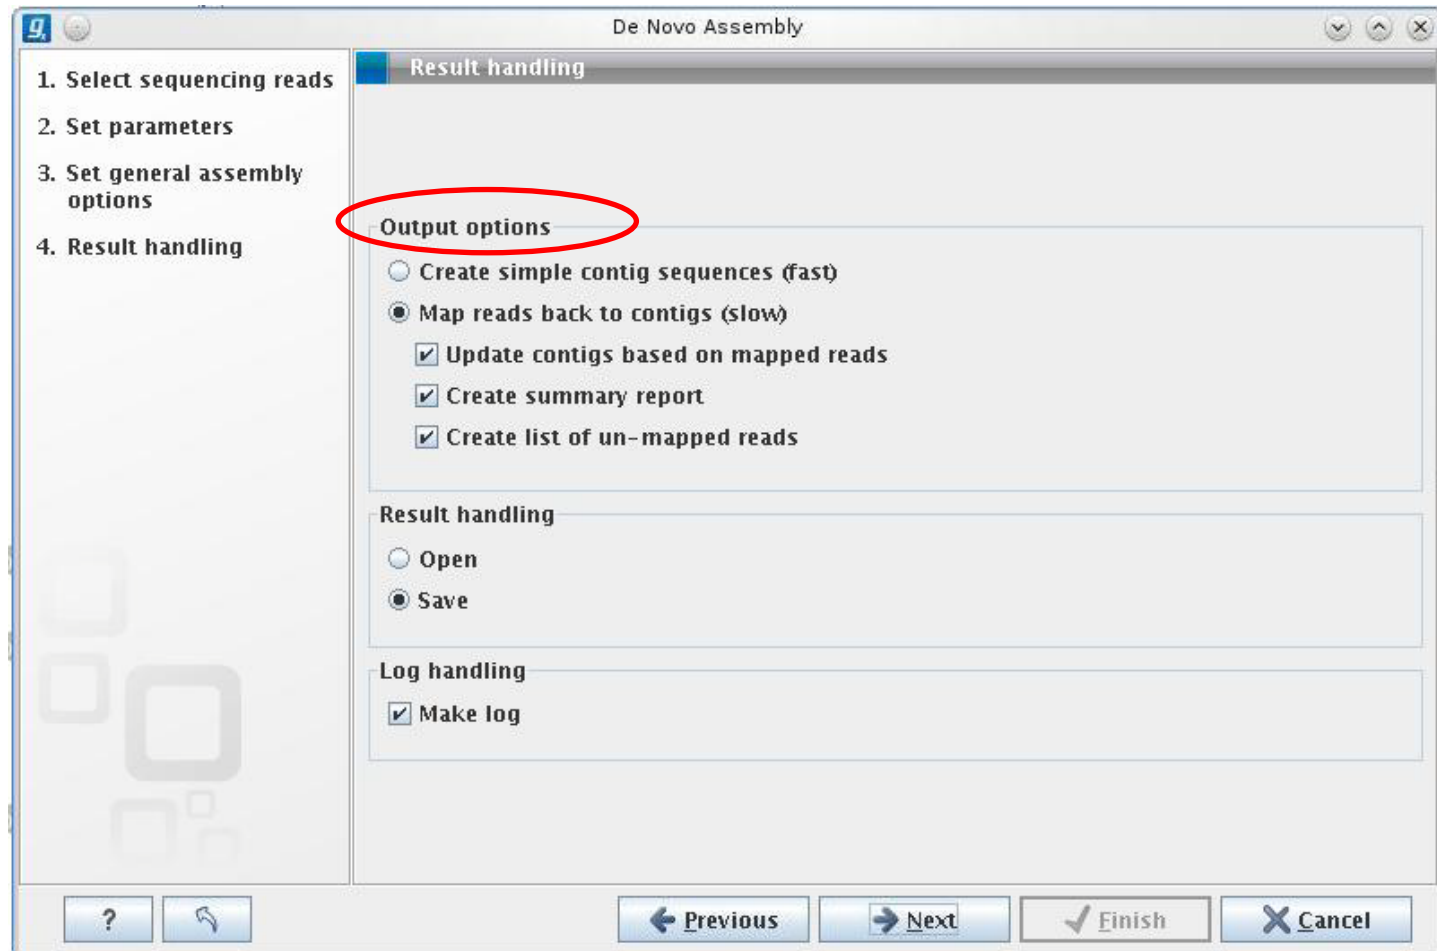
☒ Random
☐ Ignore

Contig length

Minimum contig length

? Previous Next Finish Cancel

Assembly result handling



Assembly report

Summary statistics

	Count	Average length	Total bases
Reads	629,173,407	92.43	58,156,992,907
Matched	599,285,257	92.94	55,697,623,043
Not matched	29,888,150	82.29	2,459,369,864
Contigs	2,494,829	610	1,523,965,030
Reads in pairs	162,778,034	362.64	
Broken paired reads	431,137,809	91.92	

Quality assessment

Total length of sequences (bp): 1,523,965,030

Total number of contigs: 2,494,829

Max contig length (bp): 22,049

Mean contig length (bp): 610.85

Median contig length (bp): 371

Min contig length (bp): 200

N25: 1720

N50: 891

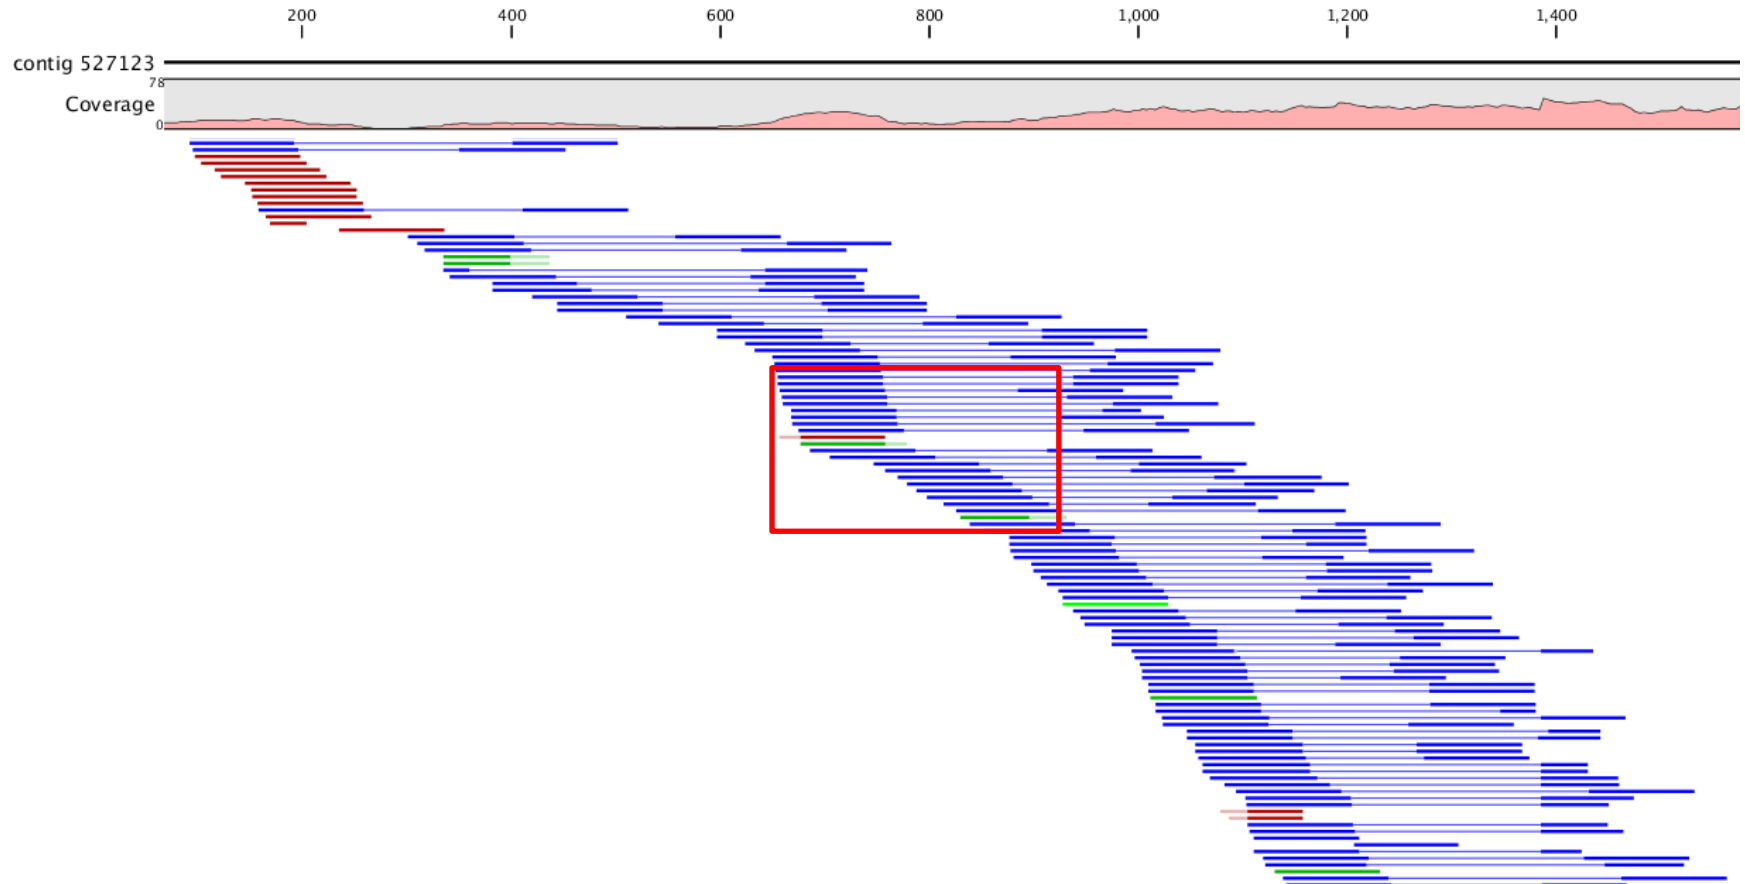
N75: 435

N90: 251

Total GC count (bp): 650,659,197

GC (%): 42.70

View assembled contig



Base level view

Coverage

Consensus sequence

Paired-end read

Single end read (Reverse)

Single end read (Forward)



Assembled contigs

Contig ID	Average coverage
>ConsensusfromContig1	Average coverage: 25.94
ACCATCATCAGCGATGATAGGGTTTACITTAGAAATTATTCGAGCAGAACTTGAAGCT AAACCTCTGAAAATTTGCAAATCCTTTTGTATCTTTTGCAGCTGAGTTTAGAATTGTACC AAAGCACTCTAATCTTTTTTTCGCAATTGATAAACTGATGCTTCTGTAGCACGTGGAGTA AACACTTCCTCTGAATTAAAGTCCTGTTGAAATGGTACAGAAACACTGCCAGAAGAACAC TTGGCACCTCAATCACAATATTCTTGTAGCACTTGGTGCGTTTA	
>ConsensusfromContig2	Average coverage: 13.80
TCCCCCTCCCCCAAAAATCAGCAACTCAAAATGGGGGTGGAGGTTTCGTGTTTCATTGCCAG GAAATACGGTAATGACGATCTAAAAGCACTCACTGAATCTTTGAGTATATTAAGGTTCTT GATTACTTGAGTTTTCTTTGCCATTCCATTAGTTCATATCTGATATGCACTTTAATTCCA ATTGTCCATTTAAAAAAAACCTATAGAGCTTATTTCTCACCTTAACCATTTTGGGGTAAA AGCAATGAGCAGTTTCTGTTTTTCAGTCATAAAATTCAGTGAGAGCTGCATAAGAAAATGT GGAGACGTCAAACATTTTTTTT	
>ConsensusfromContig3	Average coverage: 24.54
CACCCCCCCCACACACACGCACACTGTCCCGACCCTTCCCCTCACTCACGTAGAGCCGCA GCAGCAGCTCAGCTTTGACATCATTGTCAATTGCGCAGCTCGTCGGTCACAGAGTTGAAGT CCTCCTGCCATCGCGACACAACTCCTGCTTGTGATCCGGACGAATGTCCAAGTACTCCC CGTAATCCCTCCTGTGGCAGAGACACTGTCTCACTCTCCTCCCACACACACCTGTCTCAC ACCCCCACCTCATGTGTGAG	

Summary

- Genome sequencing and assembly problem
- Short read sequence assembly pipeline
 - Sequence data format (FASTQ)
 - Read quality assessment
 - Sequence trimming
 - De Novo assembly algorithms and tools
 - Assembly quality assessment
- Case study
 - Little Skate Illumina HiSeq 2000 short read sequences
 - CLC Bio Genomics Workbench

References

- Dawei Lin, Short Read Assembly, UC Davis 2010 Bioinformatics Short Course, September 13, 2010, UC Davis
- Joe Fass, Genome Sequence Assembly in Action!, UC Davis 2010 Bioinformatics Short Course, September 13, 2010, UC Davis
- NHGRI, <http://www.genome.gov/sequencingcosts/>, Feb. 4, 2011
- Phillippy AM, Schatz MC, Pop M. Genome assembly forensics: finding the elusive mis-assembly. *Genome Biol.* 2008;9(3):R55. Epub 2008 Mar 14.
- Schatz MC, Delcher AL, Salzberg SL. Assembly of large genomes using second-generation sequencing. *Genome Res.* 2010 Sep;20(9):1165-73. Epub 2010 May 27.
- CLC Bio. Manual for CLC Genomics Workbench 4.6 Windows, Mac OS X and Linux April 1, 2011
- CLC Bio. Manual for CLC Genomics Workbench 4.7 Windows, Mac OS X and Linux
- May 14, 2011
- Miller JR, Koren S, Sutton G. Assembly algorithms for next-generation sequencing data. *Genomics.* 2010 Jun;95(6):315-27. Epub 2010 Mar 6.

Thank You!

Questions???